

Titanic Pro Systems

This project is a comprehensive Machine Learning and Data Engineering application designed to predict passenger survival probabilities on the Titanic. Going beyond simple static modelling, the system integrates a full-stack data pipeline that combines historical data analysis, interactive AI predictions, and real-time streaming capabilities.

The core technology stack includes **XG Boost** for high-performance classification, **Stream lit** for the interactive web interface, and **Apache Kafka** (containerized via Docker) to handle real-time data ingestion. The application features a "Neon Deep-Ocean" themed UI, offering users three distinct modules: an Analytics Dashboard for historical insights, a Neural Survival Predictor for custom scenarios, and a Live Kafka Stream for processing passenger data in real-time.

ii. A Walkthrough of the EDA, Model Building, and Final Predictions

The project utilizes a robust machine learning pipeline defined in `train_model.py` to process raw data and generate accurate predictions.

- **Exploratory Data Analysis (EDA):** The system begins by loading the training dataset (`train.csv`). The analysis phase is visualized in the Stream lit application, which calculates key metrics such as total passenger count, overall survival rates, average fares, and age distribution.
 - **Feature Engineering & Preprocessing:** Before training, the raw data undergoes significant transformation:
 - **Title Extraction:** Passenger titles (e.g., Mr., Mrs., Dr.) are extracted from names to categorize social status.
 - **Family Grouping:** A Family Size feature is created by combining siblings (`SibSp`) and parents/children (`Parch`), alongside an Is Alone binary flag.
 - **Handling Missing Data:** The system imputes missing values using median strategies for numerical data and "most frequent" strategies for categorical data.
 - **Model Building:** The prediction engine relies on an **XG Boost Classifier** (XGB Classifier) with 100 estimators. This model is wrapped in a Scikit-Learn pipeline that automatically handles scaling (Standard Scaler) and One-Hot Encoding for categorical variables, ensuring the model is robust and production-ready.
-

iii. Demonstration of Filtering and Sorting Functionality via the Stream lit UI

The user interface acts as a dynamic control centre, allowing users to "filter" inputs and view "sorted" analytical insights through two main views in `app.py`:

- **Interactive Parameter Tuning (Filtering):** In the "Neural Survival Predictor" module, users can filter the model's input by manually adjusting specific passenger parameters. The sidebar and main expandable forms allow the selection of:
 - **Class:** 1st, 2nd, or 3rd Class.
 - **Demographics:** Gender and Age (via a slider).

- **Economic Status:** Fare price.

Upon clicking "RUN PREDICTION," the system filters these specific inputs through the trained XG Boost model to output a precise survival probability score displayed on a gauge chart.

- **Visual Data Sorting:** The "Analytics Dashboard" automatically sorts and groups historical data for quick interpretation. For example, it displays a "Survival by Class" bar chart (sorting survival rates by Pclass) and an "Age Distribution" histogram, allowing users to visually inspect how different groups fared compared to one another.
-

iv. Demonstration of the Real-Time Prediction System Using Kafka

The project simulates a modern production environment where passenger data is received continuously rather than in static batches.

- **Infrastructure:** The streaming architecture relies on **Apache Kafka** and **Zookeeper**, orchestrated via docker-compose.yml to ensure a stable message broker service.
- **The Data Producer:** A standalone script, producer.py, reads raw passenger data from test.csv and serializes it into JSON format. It acts as a live data source, pushing individual passenger records to the titanic_stream Kafka topic at one-second intervals to mimic real-time events.
- **The Streamlit Consumer:** The "Live Kafka Stream" page in the main application connects to this topic using the confluent_kafka library. As new passenger records arrive:
 1. The app deserializes the JSON data.
 2. It extracts features on the fly (e.g., calculating title and family size dynamically).
 3. The pre-trained XG Boost model predicts the survival status ("Alive" vs "Deceased") instantly.
 4. The results are logged to a live-updating table on the dashboard, showing the Passenger Name, Predicted Status, and Model Confidence score.